

Local and Emerging AI Tool Ecosystem

Local Execution Tools - Ollama: An open-source tool to run open models (Llama 3, Mistral, Gemma) locally on CPU and GPU with a simple CLI. - LM Studio: A desktop UI application to load and run GGUF format models locally, exposing an OpenAI-compatible API endpoint.

MLOps and Evaluation - LangSmith: LangChain's platform to trace, debug, test, and evaluate LLM applications and agent states. - Weights & Biases (W&B;): A tool for model tracking, hyperparameter sweeps, and logging training runs.